

Bioinformatics Project #10

Svante Sörberg (s279331)

July 2020

Abstract

Elastic Weight Consolidation (EWC) is a technique to avoid catastrophic forgetting when learning multiple tasks with a neural network. In this project, EWC has been implemented on a Convolutional Neural Network (CNN). The performance was studied on multiple tasks in the form of permutations of the MNIST dataset. The results suggest that EWC is an effective way of overcoming the problem of catastrophic forgetting for a limited amount of tasks.

Introduction

Neural networks are traditionally not well suited to multi-task learning, despite being modelled after the human brain. The key challenge is making the neural network sensitive to the requirements of the new task, while not losing the ability to perform some previously learnt task. The problem of forgetting how to perform a previous task is known as *catastrophic forgetting*.

In this project, the method of *Elastic Weight Consolidation* (EWC) proposed in [2] is studied.

Background

The principle behind EWC is to apply a kind of regularization to the weights. The most common regularization techniques penalize all weights equally, usually based on how much they differ from a set value, typically 0 or some previously trained value for each particular weight. EWC, on the other hand, penalizes weight values based on two things

1. The difference between value of the weight being evaluated and the optimal value for that weight in previous tasks
2. The *importance* of that weight for previous tasks

The idea is that weights that are important to a previously learned task should not be allowed to change very much. Weights that are inconsequential for previous tasks may change a lot to accommodate the learning of any new task.




Figure 1: An illustration of EWC compared to L2-regularization and no regularization. L2-regularization converges on a solution in parameter space that is inbetween the optimum for the respective tasks. Without regularization, the weights simply change from a good solution for task A, to a good solution for task B – completely forgetting task A in the process. With EWC, the parameters change in a direction that improves the network’s ability to solve task B, but maintaining them in a region where task A can still be solved well. The figure is sourced from [2]

The “importance” is given by the diagonal of the *Fisher Information Matrix*

(FIM). Without getting into details, the diagonal elements of the FIM for a task K , $F_{K,i}$ can be approximated (in what is known as an Empirical Fisher Matrix) by

$$F_{A,i} = \left(\frac{\partial l(\boldsymbol{\theta})}{\partial \theta_i} \right)^2 \quad \text{where } \theta_i \text{ is the } i\text{:th weight and } l(\boldsymbol{\theta}) \text{ is the loss function}$$

Clearly, if θ_i is important for task A, the gradient of the loss (for task A) will be large (positive or negative) when taken w.r.t. θ_i . The full regularization term can be written as

$$\frac{\lambda}{2} \sum_K \sum_i F_{K,i} (\theta_i - \theta_{K,i})^2 \quad \text{where } \theta_{K,i} \text{ is the learned } \theta_i \text{ for task } K$$

Note that how much emphasis is put on remembering old tasks compared to performing well on new ones can be controlled by the parameter λ .

Implementation

The implementation more or less boils down to implementing the described regularization function. At the end of the training process for each task K , the Fisher diagonal F_K as well as the final value for the trained parameters must be computed and saved, as they are needed when training subsequent tasks.

Experiments

For testing EWC, the MNIST dataset was used. To generate multiple datasets of comparable difficulty, the dataset was pixel-permuted, i.e. all pixels of the image shuffled without concern for spatial locality:



Figure 2: *To the left, random images from the MNIST dataset. To the right, random images from a pixel-permuted MNIST dataset.*

A CNN was trained, first on the regular MNIST dataset, and subsequently on four permuted datasets in order for a total of five different tasks. For details on

the network architecture and how it was trained, see appendix A. Accuracy was recorded after each training epoch on the test set of the task that was currently being trained, as well as any task that the network had already been trained on.

To demonstrate the problem of catastrophic forgetting, the first experiment trained a network on the regular MNIST dataset. The same network (without resetting weights) was then trained on all the permuted MNIST datasets in sequence. There was no penalty at all for diverging from the parameters learned for previous tasks. For the second experiment, an EWC penalty was applied for all previously learned tasks. In other words, each previously learned task was "kept in memory". For the third experiment, an EWC penalty was only applied for the previously learned task (only the previous task was kept in memory).

Results

Incremental training without penalty




Figure 3: *Clearly, this is why it is called **catastrophic** forgetting. As soon as the network is trained on a new task, performance drops significantly. While a lot better than chance, the accuracy on the first task drops from almost 100% to a mere 60% after training on the second one. On the MNIST dataset, the accuracy for random guessing would be 10%.*

Incremental training with penalty for each task




Figure 4: *Performance is high on all tasks. While the network always performs best on the most recently trained task, this does not seem to hold over longer time periods. For example, at the end, the network performs better on task 2 (orange) than on task 4 (red) which it was trained on more recently.*

Incremental training with penalty for previous task




Figure 5: *Performance on a task degrades significantly after adding two or three additional tasks. There is also a lot of fluctuation in accuracy on the older tasks.*

Discussion

Performance analysis

Using an EWC penalty associated with all previously learned tasks, rather than only the most recently learned one, clearly offers the best and most consistent performance. Interestingly though, for a limited number of tasks, a penalty for only the most recent task still offers relatively high accuracy. It is far better from an accuracy point of view than the network trained without any EWC at all.

Presumably, the reason for this is that the EWC penalty associated with remembering task i still embeds a preference for parameter values similar to those learned for task $i - 1$, even if there is no explicit reference to those values when computing the penalty. It is perhaps a little surprising that the effect is so strong, since EWC only encourages staying close to previously trained weights if they are important to that task. After all, there is no guarantee that the weights that are important for task i were also important for task $i - 1$. One might even expect the opposite to be the case: because the network is encouraged to not change the weights important for task $i - 1$, the most efficient way to train it for task i might be to more or less leave those particular weights alone unchanged and rely on other ones.

Convolutional neural network & pixel-permutations

One may question whether the pixel-permuted dataset is reasonable for evaluating a convolutional neural network. The purpose of the convolutional layers is to recognize and exploit spatial localities in the images. These local features are certainly impacted by shuffling all pixels randomly across the image. This topic is explored further, for example in [1].

It might have been more reasonable to permute the MNIST dataset in a way that somewhat preserves spatially local features of the image. One suggestion is to divide the images into (square) sub-images, e.g. dividing the 28x28 image into 16 7x7 squares and permuting the position of those squares within the image:



Figure 6: *To the left, random images from the MNIST dataset. To the right, random images from an MNIST dataset where the 7x7 sub-squares have swapped positions. Clearly this kind of permutation preserve some spatially local features such as lines, but may divide them into several pieces and/or change their position within the image. N.b. the permuted images to the right are not permutations specifically of the images to the left.*

One may even argue that in these experiments, the convolutional layer limits its performance on the permuted datasets, because the structures it learns to recognize on the original MNIST dataset (presumably lines and other simple geometric shapes) cannot be used meaningfully on the permuted ones. Perhaps, a fully connected neural network could be used for these tasks instead. Obviously, for image recognition tasks that are more complicated, CNN:s offer far better performance. It could be interesting to experiment a pre-trained model like VGG16, freezing some or all of the convolutional layers and apply EWC on the other ones, and see how well it can adapt to different image recognition tasks.

Impact on training time

Training a network with EWC is not noticeably slower than using any other regularization function. Between tasks, the fisher diagonal has to be computed and the trained parameters saved. This entails a small performance hit. The time

spent saving trained parameters and computing Fisher diagonals is proportional to the size of the network and the number of tasks trained on. While the computation of the Fisher diagonal is by no means quick, it is (for reasonably large networks) overshadowed by the time spent in actual training.

Obviously, the time spent computing the regularization penalty is proportional to the number of tasks that have been trained already, as there is one term associated with each task. It is possible to avoid this, however. Noting that:

$$\begin{aligned} & \sum_K \sum_i F_{K,i} (\theta_i - \theta_{K,i})^2 \\ &= \\ & \sum_i \left[\theta_i^2 \left(\sum_K F_{K,i} \right) - \theta_i \left(\sum_K F_{K,i} \theta_{K,i} \right) \right] + \sum_i \sum_K F_{K,i} \theta_{K,i}^2 \end{aligned}$$

it would be possible to, at the end of training for each task, compute the quantities

$$\Sigma_{1,i} = \sum_K F_{K,i} \quad , \quad \Sigma_{2,i} = \sum_K F_{K,i} \theta_{K,i} \quad \text{and} \quad \Sigma_3 = \sum_i \sum_K F_{K,i} \theta_{K,i}^2$$

for each i . Note that this simply amounts to adding the values

$$\sigma_{1,i} = F_{K^*,i} \quad , \quad \sigma_{2,i} = F_{K^*,i} \theta_{K^*,i} \quad \text{and} \quad \sigma_3 = \sum_i F_{K^*,i} \theta_{K^*,i}^2$$

computed from the most recently learned task K^* onto the quantities $\Sigma_{1,i}$, $\Sigma_{2,i}$ and Σ_3 (that are continuously updated after learning each task).

The penalty could then be computed as

$$\frac{\lambda}{2} \left(\sum_i [\theta_i^2 \cdot \Sigma_{1,i} - \theta_i \cdot \Sigma_{2,i}] + \Sigma_3 \right)$$

Which is linear in the amount of parameters instead of linear in the amount of parameters multiplied by the number of tasks learned.

References

- [1] C. Ivan. Convolutional neural networks on randomized data. In *CVPR Workshops*, pages 1–8, 2019.
- [2] J. Kirkpatrick, R. Pascanu, N. Rabinowitz, J. Veness, G. Desjardins, A. A. Rusu, K. Milan, J. Quan, T. Ramalho, A. Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13):3521–3526, 2017.

A Network architecture & parameters

For all experiments, the following Keras sequential model was used:

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(28,28,1), name='input'),
    tf.keras.layers.Dropout(rate=0.2, name='dropout1'),
    tf.keras.layers.Conv2D(
        64,
        (3, 3),
        padding='same',
        activation='relu',
        name='conv1-1'
    ),
    tf.keras.layers.MaxPooling2D(
        pool_size=(2,2),
        name='pool1'
    ),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(
        1024,
        activation='relu',
        name='fc1'
    ),
    tf.keras.layers.Dropout(rate=0.5, name='dropout2'),
    tf.keras.layers.Dense(
        1024,
        activation='relu',
        name='fc2'
    ),
    tf.keras.layers.Dropout(rate=0.5, name='dropout3'),
    tf.keras.layers.Dense(
        10,
        activation='softmax',
        name='output'
    )
])
```

Training parameters were:

Training epochs (per task)	30
Batch size	64
Learning rate	10^{-4}
λ used for EWC	200
Keras optimizer	ADAM
Keras loss function	Categorical Crossentropy